




# Model description: GPT2

GPT-2 is a transformers model pretrained on a very large corpus of English data in a self-supervised fashion. This means it was pretrained on the raw texts only, with no humans labelling them in any way (which is why it can use lots of publicly available data) with an automatic process to generate inputs and labels from those texts. More precisely, it was trained to guess the next word in sentences. we use `small` gpt model.



This is relevant for models that support tasks involving multiple inputs

## Input:

This is relevant for models that support tasks involving multiple inputs

- 1) Tokenization.
- 2) Special Tokens trained for german language.
- 3) Padding.
- 4) Segmentation.
- 5) Conversion to Input IDs.

## Output:

The output of a GPT-2 model is typically a sequence of tokens representing the generated text. These tokens are numerical IDs that correspond to the indices of the tokens in the model's vocabulary. The output tokens can be decoded back into human-readable text using the model's tokenizer

# Steps to make model text\_generation

## 1) Data Collection: "German Recipes Dataset"

data set is german language contain: **12190 german recipes**

### content:

- Ingredients: the ingredients of the recipe as array
- Instructions: the instructions as free text
- Name: the name of the recipe
- Url: the source url
- Day: the day where the recipe was created
- Month: the month where the recipe was created
- Year: the year where the recipe was created
- Weekday: the weekday where the recipe was created

**2) Preprocessing:** Clean the data by removing irrelevant characters, formatting issues, or any other noise. This step might also include tokenization (splitting the text into words or subwords), lowercasing, and removing stopwords.

**3) Tokenization:** return dictionary contain( `input_types, input_ids, attention_mask` )

### From tokens to input IDs

The conversion to input IDs is handled by the `convert_tokens_to_ids()` tokenizer method:

### Attention masks

*Attention masks* are tensors with the exact same shape as the input IDs `tensor`, filled with `0s` and `1s`: `1s` indicate the corresponding tokens should be attended to, and `0s` indicate the corresponding tokens should not be attended to (i.e., they should be ignored by the attention layers of the model).

we use transformers "Pytorch"

```

from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("anonymous-german-nlp,

train_path = 'train_dataset.txt'
test_path = 'test_dataset.txt'

```

**4) Model Selection:** we use model **"German\_GPT2"** [dbmdz/german-gpt2](#)

**5) Training:** we used pretrained model ["anonymous-german-nlp/german-gpt2"](#)  
we split data in (train and test)

**6) Fine-tuning:** Fine-tuning the model allows users to customize the pre-trained GPT-2 model to suit their specific needs, improving its overall performance and utility in niche domains.

steps for fine-tuning: 1-inastall **"transformers"** and use **pytorch**

2-Importing Necessary Libraries such as  
(Trainer, TrainingArguments, AutoModelWithLMHead, TextDataset,  
DataCollatorForLanguageModeling)

3-

**a)**Loading the GPT-2 Model and Tokenizer. **b)**Loading the Training Dataset.

**c).** Creating Data Collator. **d)** Setting Training Arguments.

**e)** Initialising the Trainer. **f)** Training the Model. **g)**Saving the Fine-Tuned Model

## Training the Model with Text Data

```

from transformers import Trainer, TrainingArguments, AutoModelWithLMHead

model = AutoModelWithLMHead.from_pretrained("anonymous-german-nlp,

training_args = TrainingArguments(
    output_dir="folder", #The output directory
    overwrite_output_dir=True, #overwrite the content of the output directory
    num_train_epochs=2, # number of training epochs

```

```

per_device_train_batch_size=32, # batch size for training
per_device_eval_batch_size=64, # batch size for evaluation
eval_steps = 400, # Number of update steps between two evals
save_steps=800, # after # steps model is saved
warmup_steps=500, # number of warmup steps for learning rate
prediction_loss_only=True,
report_to="tensorboard"
)

```

9) **Generation:** Once the model is trained, it can generate text by providing a prompt or seed input. The model then generates the subsequent text based on its learned patterns and probabilities.

```

from transformers import pipeline

```

```

generation = pipeline('text-generation', model='anonymous-german

```

## Internal structure for attention for text generation

### The Encoder-Decoder Framework

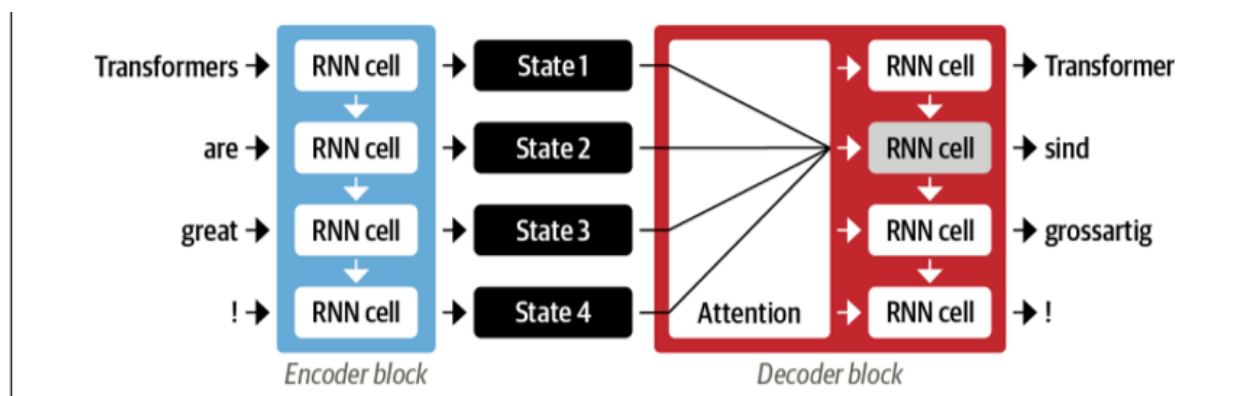


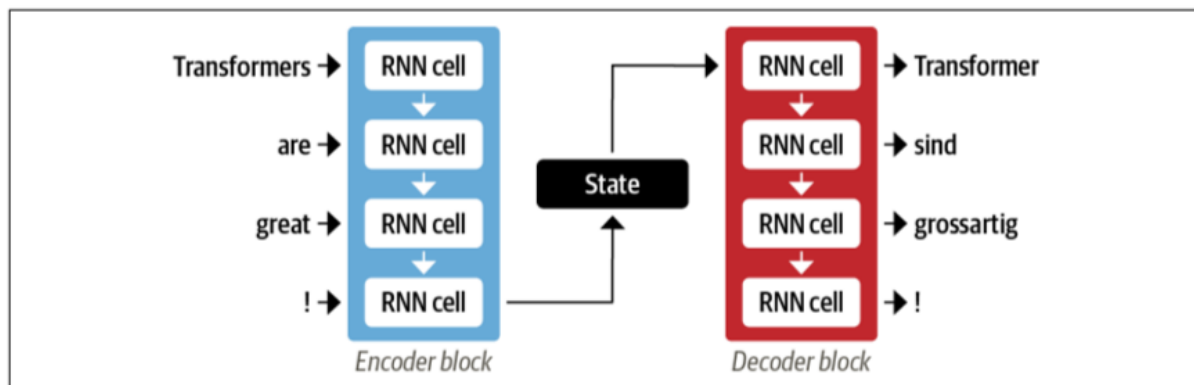
Figure 1-4. An encoder-decoder architecture with an attention mechanism for a pair of RNNs

## Decoding

*Decoding* is going the other way around: from vocabulary indices, we want to get a string.

This can be done with the `decode()` method as follows:

```
decoded_string = tokenizer.decode([7993, 170, 11303, 1200, 24
43, 1110, 3014])
print(decoded_string)
```



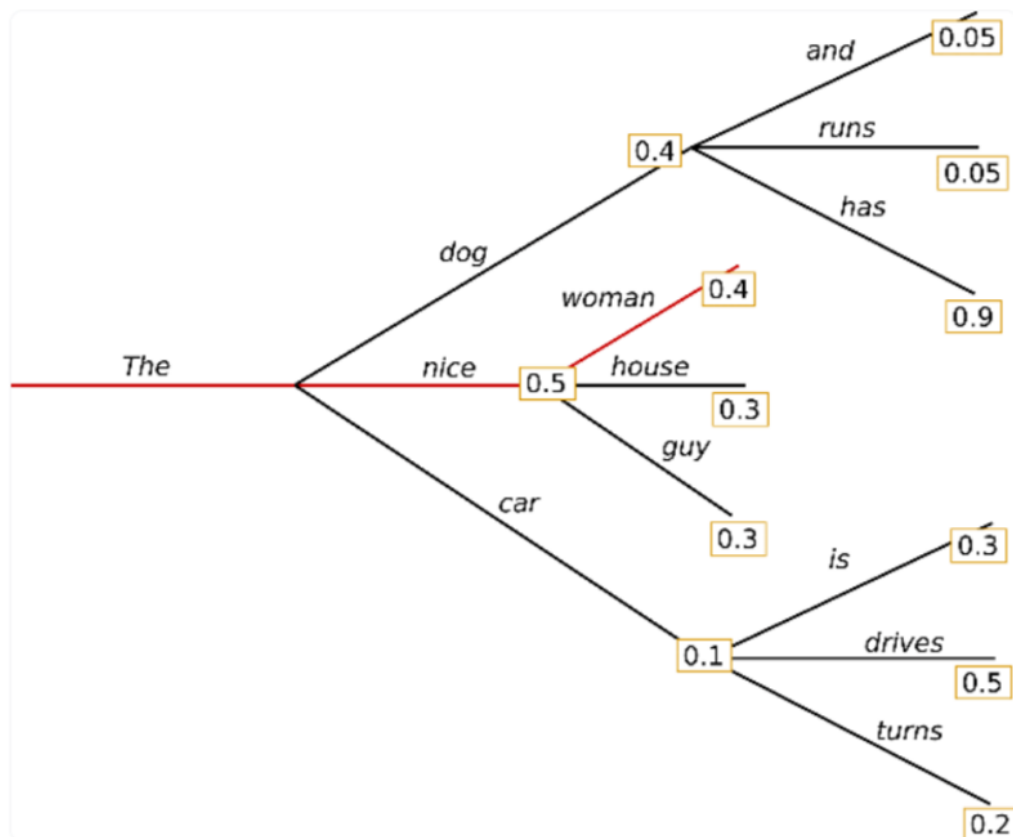
*Figure 1-3. An encoder-decoder architecture with a pair of RNNs (in general, there are many more recurrent layers than those shown here)*

## Greedy Search

`generate` uses greedy search decoding by default so you don't have to pass any parameters to enable it. This means the parameters `num_beams` is set to 1 and `do_sample=False`.

## Greedy Search

Greedy search is the simplest decoding method. It selects the word with the highest probability as its next word:  $w_t = \operatorname{argmax}_w P(w|w_{1:t-1})$  at each timestep  $t$ . The following sketch shows greedy search.



Starting from the word "The", the algorithm greedily chooses the next word of highest

## Masked self\_Attention

An attention mask is a **binary mask that designates which tokens should be attended to (assigned non-zero weights) and which should be ignored (assigned zero weights)**

such as (Full Attention, Scaled Dot-Product Attention, Masked Attention)

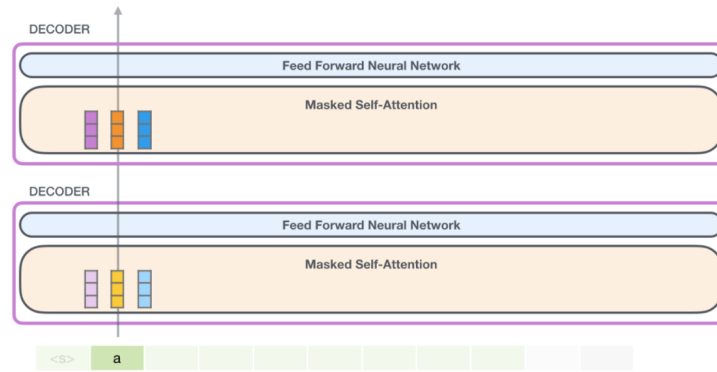
## GPT-2 Masked Self-Attention

Let's get into more detail on GPT-2's masked attention.

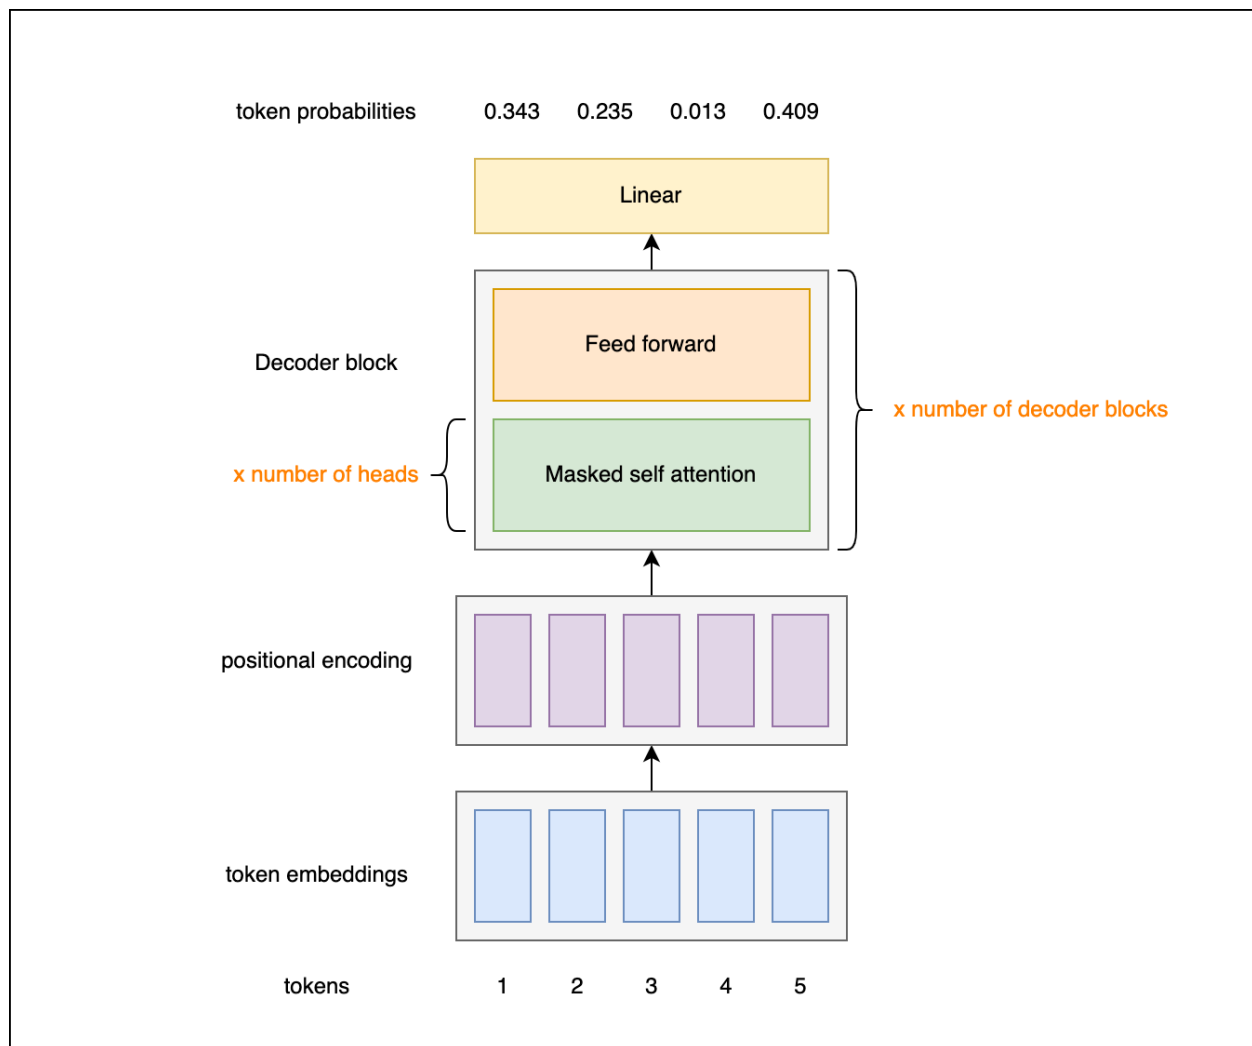
### Evaluation Time: Processing One Token at a Time

We can make the GPT-2 operate exactly as masked self-attention works. But during evaluation, when our model is only adding one new word after each iteration, it would be inefficient to recalculate self-attention along earlier paths for tokens that have already been processed.

In this case, we process the first token (ignoring `<S>` for now).







## Comparison between models :

### 1) GPT-2 Model:

#### Description:

The GPT-2 (Generative Pre-trained Transformer 2) model is a type of language model developed by OpenAI. It is based on the transformer architecture and is trained on a large corpus of text data to generate human-like text. The model is capable of generating coherent and contextually relevant text given a prompt. It does this by predicting the next word in a sequence of text based on the words

that have come before it. and it's `##usage` is instantiated the same way as in the Transformers library. The only difference is that there are a few new training arguments specific to HPUs.

## Definition:

**Output Hidden States:** If set to True, the model will return all hidden states. This can be useful for tasks like fine-tuning or analysis.

**Output Attentions:** If set to True, the model will return attention weights. This is helpful for understanding which parts of the input the model pays attention to.

**Output Cross-Attention:** If set to True, the model will return cross-attention weights. This is relevant for models that support tasks involving multiple inputs.

**Num Layers:** Specifies the number of transformer layers in the model.

**Num Heads:** Specifies the number of attention heads in each transformer layer.

**Hidden Size:** Specifies the size of the hidden layers in the transformer network.

**Intermediate Size:** Specifies the size of the intermediate (feed-forward) layers in the transformer network.

**Activation Function:** Specifies the activation function used in the feed-forward layers, typically "gelu" or "relu".

**Dropout Rate:** Specifies the dropout probability for dropout layers within the model.

**Attention Dropout Rate:** Specifies the dropout probability for attention scores in the attention layers.

**Initializer Range:** Specifies the range for random weight initialization.

These are just some of the parameters available for configuring the GPT-2 model in Hugging Face's transformers library.

Depending on your specific use case, you may need to adjust these parameters to achieve optimal performance or to suit your task requirements.

and ,

### **Model Size:**

GPT-2 comes in different sizes, ranging from "small" (117M parameters) to "large" (1.5B parameters). You can specify the model size using the `model_name_or_path` parameter when initializing the model.

### **Context length:**

The GPT-2 model has a context length of 1024 tokens for its base version and up to 8192 tokens for larger versions.

## Advantages:

- 1-strength: provides "true" features for text generation and can be used as a flexible framework for learning text representations
- 2-Accuracy: GPT-2 effectively detects machine-generated text with high accuracy
- 3-Empowering Label Mapping: GPT-2 has strong generative power and can enhance semantic embedding and improve mapping to relevant labels.
- 4-Community Support: Hugging Face has a large community of users and contributors, providing support and resources for working with GPT-2 and other models.
- 5-High-quality text generation: With a wide range of uses GPT-2 is known for its capacity to produce high-quality human-like writing.
- 6-Pre-trained models: The pre-trained models included with GPT-2 may be utilized for a variety of applications involving Natural Language Processing, without the need for an additional training process.
- 7-Large-scale architecture: The architecture of GPT-2 is built to handle enormous volumes of data, making it appropriate for applications that need to analyze massive datasets.
- 8-Flexibility: GPT-2 is tailored to perform a range of Natural Languages Processing tasks, such as question answering, text summarization, and language translation, and it can be fine-tuned on specific data sets or tasks, making it adaptable to different applications and domains, provides "true" features for text generation and can be used as a flexible framework for learning text representations
- 9-Ease of Use: GPT-2 is relatively easy to use compared to some other AI models,

requiring minimal input to generate text.

10-Open Source: The model and its codebase are open source, enabling researchers and developers to use and build upon it freely.



## Disadvantages:

1-Computational Resources: GPT-2 is a large model that requires significant

computational resources for training and inference, which can be costly and

time-consuming, especially for large-scale applications.

2-Fine-tuning Complexity: Fine-tuning the model for specific tasks or datasets can be

complex and may require expertise in natural language processing (NLP) and machine learning.

3-Limited Context Understanding: While GPT-2 can generate coherent text, it may

sometimes struggle with understanding complex contexts or generating long-range dependencies in text.

4-Ethical Considerations: GPT-2 has raised concerns about the potential misuse of AI-generated text for spreading misinformation or generating harmful content, leading to ethical considerations regarding its deployment.

5-Model Size: The size of the GPT-2 model can be a disadvantage in some applications where resource constraints are a concern, as it may not be feasible to deploy the model in environments with limited resources.

6-Lack of Control: While GPT-2 can generate high-quality text, it may sometimes produce outputs that are inappropriate or unintended, requiring careful monitoring and control

mechanisms.

7-Bias: Like any model trained on human-generated data, GPT-2 can exhibit biases present in the training data.

## 2) Mistral-7B-Instruct-v0.2 Model:

### Description:

The Mistral-7B-Instruct-v0.2 model is likely a variant of a natural language processing model developed for text generation tasks. It could be based on architectures like Transformers or recurrent neural networks (RNNs). The "7B" might refer to the number of parameters in the model, indicating its size and complexity. The "Instruct" part of the name might suggest a focus on generating instructional or procedural text.

### Definition :

has the same definition of GPT-2 Model

|

but the difference is in the model size

### Model Size:

Different variants of Mistral may have different sizes or numbers of parameters. The model size can impact its performance and computational requirements.

### Context length:

The context length for the Mistral-7B-Instruct-v0.2 model of text generation is 32,000 tokens , This context window is significantly larger than the 8,000-token context window of its predecessor, **Mistral-7B-v0.1**

## Advantages:

- 1-Enhanced context: The larger context window enables the model to use more context-aware information, leading to improved responses.
- 2-Fine-tuning flexibility: Developers can easily fine-tune this model for specific tasks and applications, making it versatile and adaptable.
- 3-Large model size (7 billion parameters): A larger model typically captures more linguistic nuances, semantics, and context, leading to more coherent and accurate text generation compared to smaller models.
- 4-Adaptability: The model might be fine-tuned on specific instruction sets or datasets for particular tasks, enhancing its performance and relevance to particular domains.

## Disadvantages:

- 1-Lack of moderation mechanisms: The model does not include any built-in moderation features, which could be a limitation in certain use cases.
- 2-Computational requirements: A model with 7 billion parameters requires substantial computational resources for training and inference, making it less accessible to users without access to powerful hardware.
- 3-Data biases: Like other large language models, Mistral-7B-Instruct-v0.1 may have biases present in the training data, potentially leading to biased or inappropriate instructional outputs.
- 4-Lack of interpretability: Large models like Mistral-7B-Instruct-v0.1 are often described as "black boxes," making it challenging to understand the internal decision-making process or troubleshoot errors.

## 3)Google/gemma-2b Model:

## Description:

Gemma is a family of lightweight, state-of-the-art open models from Google, built from the same research and technology used to create the Gemini models. They are text-to-text, decoder-only large language models, available in English, with open weights, pre-trained variants, and instruction-tuned variants. Gemma models are well-suited for a variety of text generation tasks, including question answering, summarization, and reasoning. Their relatively small size makes it possible to deploy them in environments with limited resources such as a laptop, desktop or your own cloud infrastructure, democratizing access to state of the art AI models and helping foster innovation for everyone.

## Definition :

has the same definition of GPT-2 Model

|

but the difference is in the model size

## Model Size:

Different variants of Mistral may have different sizes or numbers of parameters. The model size can impact its performance and computational requirements.

## Context length:

Gemma 2B models are trained on a context length of 8192 tokens.

## Advantages:

**1-Relatively Small Size:** Compared to other large language models, Gemma is relatively small. This makes it easier to deploy on devices with limited resources, such as laptops, desktops, or even personal cloud infrastructure. This wider accessibility fosters innovation and democratizes access to advanced AI for text generation.



**2-State-of-the-Art:** Despite their smaller size, they offer competition to their efficient architecture.

**3-Democratizing Access:** By being accessible to a wider audience, and creativity.

**4-Open-Source Availability:** Gemma-2b's open-source nature allows transparency, customization, and collaboration within the AI community. Access to the underlying code and weights, fine-tune the model for specific tasks, and contribute to its development.

## Disadvantages:

**1-Limited Capabilities:** Gemma-2b might not be as powerful as some other models on the market. This could show up in areas like handling complex tasks or generating highly nuanced creative text formats.

**2-Potential for Biases:** Like many AI models, Gemma-2b can inherit biases from the data it's trained on. This can lead to outputs that are discriminatory or biased. It's crucial to be mindful of these potential biases and implement safeguards.

**3-Fine-Tuning Expertise Required:** To unlock the full potential of Gemma-2b for specific tasks, some expertise in fine-tuning large language models might be necessary. This could limit accessibility for users who are not familiar with AI or machine learning.

<https://huggingface.co/google-bert/bert-base-cased>